

CPSC 340: Machine Learning and Data Mining

Finding Similar Items


Fall 2017

Admin

- **Assignment 1** is due tonight.
 - 1 late day to hand in Monday, 2 late days for Wednesday.
- **Assignment 2** will be up soon.
 - Start early.
- We'll start using **gradients and linear algebra** next week:
 - Many people get lost when we get to this material.
 - If you aren't comfortable with these, start reviewing/practicing!

Last Time: Outlier Detection

- We discussed **outlier detection**:
 - Identifying “unusually” different objects.
 - **Hard to precisely define.**
- We discussed 3 common approaches:
 - **Fit a model**, see if points fit the model.
 - **Plot the data**, and look for weird points.
 - **Cluster the data**, and see if points don't cluster.



Distance-Based Outlier Detection


- Most outlier detection approaches are **based on distances**.
- Can we skip model/plot/clustering and **just measure distances**?
 - How many points lie in a radius 'r'?
 - What is distance to k^{th} nearest neighbour?
- UBC connection (first paper on this topic):

Algorithms for Mining Distance-Based Outliers in Large Datasets

Edwin M. Knorr and Raymond T. Ng
Department of Computer Science
University of British Columbia


Global Distance-Based Outlier Detection: KNN

- KNN outlier detection:
 - For each point, compute the **average distance to its KNN**.
 - Sort the set of 'n' average distances.
 - Choose the biggest values as outliers.
 - **Filter out points that are far from their KNNs.**
- Goldstein and Uchida [2016]:
 - Compared 19 methods on 10 datasets.
 - **KNN best for finding “global” outliers.**
 - “Local” outliers best found with **local distance-based** methods...



Local Distance-Based Outlier Detection



- As with density-based clustering, **problem with differing densities:**



- Outlier o_2 has similar density as elements of cluster C_1 .
- Basic idea behind **local distance-based** methods:
 - Outlier o_2 is “**relatively**” **far** compared to its neighbours.


Local Distance-Based Outlier Detection

- “Outlierness” ratio of example ‘i’:
$$\frac{\text{average distance of 'i' to its } KNN_5}{\text{average distance of neighbours of 'i' to their } KNN_5}$$
- If outlierness > 1 , x_i is further away from neighbours than expected.



Problem with Unsupervised Outlier Detection

- Why wasn't the hole in the ozone layer discovered for 9 years?



- Can be **hard to decide when to report** an outlier:
 - If **you report too many non-outliers, users will turn you off.**
 - Most antivirus programs do not use ML methods (see ["base-rate fallacy"](https://en.wikipedia.org/wiki/Base_rate_fallacy))


Supervised Outlier Detection

- Final approach to outlier detection is to use supervised learning:
 - $y_i = 1$ if x_i is an outlier.
 - $y_i = 0$ if x_i is a regular point.
- We can use our methods for supervised learning:
 - We can find very complicated outlier patterns.
 - Classic credit card fraud detection methods used decision trees.
- But it needs supervision:
 - We need to know what outliers look like.
 - We may not detect new “types” of outliers.

(pause)

Motivation: Product Recommendation

- A customer comes to your website looking to buy at item:



- You want to **find similar items** that they might also buy:


Customers Who Bought This Item Also Bought

Page 1 of 20



User-Product Matrix


Column gives
all users that
bought product.



Row x_i gives all items bought by user 'i'. By convention, x_i is a $d \times 1$ column vector.

Amazon Product Recommendation

- Amazon product recommendation method:



- Return the **KNNs across columns**.
 - Find 'j' values minimizing $||x^i - x^j||$.
 - **Products that were bought by similar users.**
- But first **divide each column by its norm**, $x^i / ||x^i||$.
 - This is called **normalization**.
 - Reflects whether product is bought by many people or few people.

Amazon Product Recommendation

- Consider this user-item matrix:

$X =$

	Product 1	Product 2	Product 3	Product 4	Product 5	Product 6
John	1	1	1	1	0	1
Paul	1	0	1	0	1	0
George	1	0	1	0	1	1
Ringo	1	0	1	0	1	1
Yoko	1	1	0	1	0	0

- Product 1 is most similar to Product 3 (bought by lots of people).
- Product 2 is most similar to Product 4 (also bought by John and Yoko).
- Product 3 is **equally similar to Products 1, 5, and 6**.
 - Does not take into account that Product 1 is more popular than 5 and 6.

Amazon Product Recommendation

- Consider this user-item matrix (**normalized**):

$X =$

	Product 1	Product 2	Product 3	Product 4	Product 5	Product 6
John	$1/\sqrt{5}$	$1/\sqrt{2}$	$1/\sqrt{4}$	$1/\sqrt{2}$	0	$1/\sqrt{3}$
Paul	$1/\sqrt{5}$	0	$1/\sqrt{4}$	0	$1/\sqrt{3}$	0
George	$1/\sqrt{5}$	0	$1/\sqrt{4}$	0	$1/\sqrt{3}$	$1/\sqrt{3}$
Ringo	$1/\sqrt{5}$	0	$1/\sqrt{4}$	0	$1/\sqrt{3}$	$1/\sqrt{3}$
Yoko	$1/\sqrt{5}$	$1/\sqrt{2}$	0	$1/\sqrt{2}$	0	0

- Product 1 is most similar to Product 3 (bought by lots of people).
- Product 2 is most similar to Product 4 (also bought by John and Yoko).
- Product 3 is **most similar to Product 1**.
 - Normalization means it **prefers the popular items**.

Cost of Finding Nearest Neighbours

- With 'n' users and 'd' products, finding KNNs costs $O(nd)$.
 - Not feasible if 'n' and 'd' are in the millions.
- It's faster if the user-product matrix is sparse: $O(z)$ for z non-zeroes.
 - But 'z' is still enormous in the Amazon example.

Closest-Point Problems

- We've seen a lot of “closest point” problems:
 - K-nearest neighbours classification.
 - K-means clustering.
 - Density-based clustering.
 - Hierarchical clustering.
 - KNN-based outlier detection.
 - Outlierness ratio.
 - Amazon product recommendation.
- How can we possibly apply these to Amazon-sized datasets?

But first the easy case: “Memorize the Answers”

- Easy case: you have a **limited number of possible test examples**.
 - E.g., you will always choose an existing product (not arbitrary features).
- In this case, just **memorize the answers**:
 - For each test example, compute all KNNs and store pointers to answers.
 - At test time, just return a set of pointers to the answers.
- The answers are called an **inverted index**, queries now cost $O(k)$.
 - Needs an extra $O(nk)$ storage.


Grid-Based Pruning

- Assume we want to find objects within a distance of 'r' of point x_i .

Divide space
into squares
of length r .


Hash examples based on
squares:

$\text{Hash}["64,76"] = \{x_3, x_{70}\}$
(Dict in Python/Julia)



Grid-Based Pruning

- Which squares do we need to check?




Points in **same square** can have distance less than 'r'.

Grid-Based Pruning

- Which squares do we need to check?


Points in **adjacent squares** can have distance less than distance 'r'.



Grid-Based Pruning

- Which squares do we need to check?

Points in **non-adjacent squares** must have distance **more than 'r'**.



Grid-Based Pruning


- Assume we want to find objects within a distance of ' r ' of point x_i .

Divide space
into squares
of length r .

Hash examples based on
squares:




$\text{Hash}["64,76"] = \{x_3, x_{70}\}$
(Dict in Python/Julia)

Only need to check
points in same and
adjacent squares.



Grid-Based Pruning Discussion

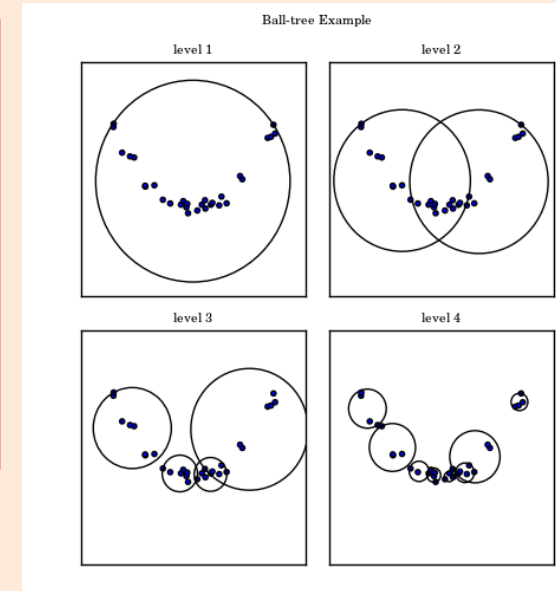


- Similar ideas can be used for other “closest point” calculations.
 - Can be used with any norm.
 - If you want KNN, can use need grids of multiple sizes.
- But we have the “curse of dimensionality”:
 - Number of adjacent regions increases exponentially:
 - 2 with $d=1$, 8 with $d=2$, 26 with $d=3$, 80 with $d=4$, 252 with $d=5$, $3^d - 1$ in d -dimension.



Grid-Based Pruning Discussion

- Better choices of regions:

- Quad-trees.
- Kd-trees.
- R-trees.
- Ball-trees.




- Works better than squares, but **worst case is still exponential.**

Approximate Nearest Neighbours

- *Approximate* nearest neighbours:
 - We allow errors in the nearest neighbour calculation to gain speed.
- A simple and very-fast approximate nearest neighbour method:
 - Only check points within the same square.
 - Works if neighbours are in the same square.
 - But misses neighbours in adjacent squares.
- A simple trick to improve the approximation quality:
 - Use more than one grid.
 - So “close” points have more “chances” to be in the same square.

Approximate Nearest Neighbours


Grid 1:



Approximate Nearest Neighbours


- Using multiple sets of regions improves accuracy.

Grid 2:



Approximate Nearest Neighbours

- Using multiple sets of regions improves accuracy.



Locality-Sensitive Hashing

- Even with multiple regions, **approximation can be poor for large 'd'**.
- Common Solution (**locality-sensitive hashing**):
 - Replace features x_i with lower-dimensional features z_i .
 - E.g., turns each a 1000000-dimensional x_i into a 10-dimensional z_i .
 - Choose **random z_i to preserve high-dimensional distances** (bonus slides).

$$\|z_i - z_j\| \approx \|x_i - x_j\|$$

- Find points hashed to the same square in lower-dimensional ' z_i ' space.
- **Repeat** with different random z_i values to increase chances of success.

End of Part 2: Key Concepts

- We focused on 3 unsupervised learning tasks:
 - Clustering.
 - Partitioning (k-means) vs. density-based.
 - “Flat” vs. hierarachial (agglomerative).
 - Vector quantization.
 - Label switching.
 - Outlier Detection.
 - Ambiguous objective.
 - Common approaches (model-based, graphical, clustering, distance-based, supervised).
 - Finding similar items.
 - Amazon product recommendation.
 - Region-based pruning for fast “closest point” calculations.
- If previous years we also covered “association rules”:
 - <http://www.cs.ubc.ca/~schmidtm/Courses/340-F16/L12.pdf>

Summary

- Distance-based outlier detection:
 - Based on measuring (relative) distance to neighbours.
- Supervised-learning for outlier detection:
 - Can detect complex outliers given a training set.
- Amazon product recommendation:
 - Find similar items using nearest neighbour search.
- Fast nearest neighbour methods drastically reduce search time.
 - Inverted indices, distance-based pruning.
- Next week: how do we do supervised learning with a *continuous* y_i ?

Locality-Sensitive Hashing

- How do we make distance-preserving low-dimensional features?
- Johnson-Lindenstrauss lemma (paraphrased):
 - Define element 'j' of 'z_i' by:
$$z_{ij} = w_{j1} x_{i1} + w_{j2} x_{i2} + \dots + w_{jd} x_{id}$$
 - Where the scalars 'w_{jc}' are samples from a standard normal distribution.
 - We can collect them into a matrix 'W', which is the same for all 'i'.
 - If the dimension 'k' of the 'z_i' is large enough, then: $\|z_i - z_j\| \approx \|x_i - x_j\|$
 - Specifically, we'll require $k = \Omega(\log(d))$.

Locality-Sensitive Hashing

- Locality-sensitive hashing:
 1. Multiply X by a random Gaussian matrix 'W' to reduce dimensionality.
 2. Hash dimension-reduced points into regions.
 3. Test points in the same region as potential nearest neighbours.
- Now repeat with a different random matrix.
 - To increase the chances that the closest points are hashed together.
- An accessible overview is here:
 - <http://www.slaney.org/malcolm/yahoo/Slaney2008-LSHTutorial.pdf>

Cosine Similarity vs. Normalized Nearest Neighbours

- The Amazon paper says they “maximize cosine similarity”.
- But this is equivalent to **normalized nearest neighbours**.
- Proof for $k=1$:

$$\begin{aligned} \underbrace{\operatorname{argmin}_j \left\| \frac{x_i}{\|x_i\|} - \frac{x_j}{\|x_j\|} \right\|}_{\text{normalized nearest neighbour}} &\equiv \operatorname{argmin}_j \frac{1}{2} \left\| \frac{x_i}{\|x_i\|} - \frac{x_j}{\|x_j\|} \right\|^2 \\ &\equiv \operatorname{argmin}_j \frac{1}{2} \frac{x_i^T x_i}{\|x_i\|^2} - \frac{2 x_i^T x_j}{\|x_i\| \cdot \|x_j\|} + \frac{1}{2} \frac{x_j^T x_j}{\|x_j\|^2} \\ &\equiv \operatorname{argmin}_j - \frac{x_i^T x_j}{\|x_i\| \cdot \|x_j\|} \\ &\equiv \operatorname{argmax}_j \frac{x_i^T x_j}{\|x_i\| \cdot \|x_j\|} \quad \left. \vphantom{\operatorname{argmax}_j} \right\} \rightarrow \text{maximum cosine similarity} \end{aligned}$$

Outlierness (Symbol Definition)

- Let $N_k(x_i)$ be the **k-nearest neighbours** of x_i .
- Let $D_k(x_i)$ be the **average distance** to k-nearest neighbours:

$$D_k(x_i) = \frac{1}{k} \sum_{j \in N_k(x_i)} \|x_i - x_j\|$$


- **Outlierness** is ratio of $D_k(x_i)$ to average $D_k(x_j)$ for its neighbours 'j':

$$O_k(x_i) = \frac{D_k(x_i)}{\frac{1}{k} \sum_{j \in N_k(x_i)} D_k(x_j)}$$

- If outlierness > 1 , x_i is **further away from neighbours** than expected.

Outlierness with Close Clusters


- If clusters are close, outlierness gives unintuitive results:



- In this example, 'p' has higher outlierness than 'q' and 'r':
 - The green points are not part of the KNN list of 'p' for small 'k'.

Outlierness with Close Clusters

- ‘Influenced outlierness’ (INFLO) ratio:
 - Include in denominator the ‘reverse’ k-nearest neighbours:
 - Points that have ‘p’ in KNN list.
 - Adds ‘s’ and ‘t’ from bigger cluster that includes ‘p’:



- But still has problems:
 - Dealing with hierarchical clusters.
 - Yields many false positives if you have “global” outliers.
 - Goldstein and Uchida [2016] recommend just using KNN.

Malware and Intrusion Detection Systems

- In antivirus software and software for network intrusion detection systems, another method of outlier detection is common:
 - “Signature-based” methods: keep a list of byte sequences that are known to be malicious. Raise an alarm if you detect one.
 - Typically looks for **exact** matches, so can be implemented very quickly.
 - E.g., using data structures like “suffix trees”.
 - Can’t detect new types of outliers, but if you are good at keeping your list of possible malicious sequences up to date then this is very effective.
 - Here is an article discussing why ML is **not** common in these settings:
 - <http://www.icir.org/robin/papers/oakland10-ml.pdf>

Shingling: Decomposing Objects into Pars

- We say that a program is a virus if it has a malicious byte sequence.
 - We **don't try to compute similarity of the whole program**.
- This idea of finding similar “parts” is used in various places.
- A key tool to help us do this is “**shingling**”:
 - Dividing an object into consecutive “parts”.
 - For example, we previously saw “bag of words”.
- Given the shingles, we can **search for similar parts** rather than whole objects.

Shingling Applications

- For example, n-grams are one way to shingle text data.
 - If we use tri-grams, the sentence “there are lots of applications of nearest neighbours” would have these shingles:
 - {“there are lots”, “are lots of”, “lots of applications”, “of applications of”, “applications of nearest”, “of nearest neighbours”}.
 - We can find similar items using similarity/distance between sets.
 - For example, using the Jaccard similarity.
- Applications where **finding similar shingles** is useful:
 - Detecting plagiarism (shared n-grams indicates copying).
 - BLAST gene search tool (shingle parts of a biological sequence).
 - Entity resolution (finding whether two citations refer to the same thing).
 - Fingerprint recognition (shingles are “minutiae” in different image grid cells).

Shingling Practical Issues

- In practice, you can save memory by not storing the full shingles.
- Instead, define a **hash function mapping from shingles to bit-vectors**, and just store the bit-vectors.
- However, for some applications even storing the bit-vectors is too costly:
 - This leads to randomized algorithms for computing Jaccard score between huge sets even if you don't store all the shingles.
- Conceptually, it's still useful to think of the “bag of shingles” matrix:
 - X_{ij} is '1' if object 'i' has shingle 'j'.

Minhash and Jaccard Similarity

- Let $h(x_i)$ be the smallest index 'j' where x_{ij} is non-zero ("minhash").
- Consider a **random permutation** of the possible shingles 'j':
 - In Julia: `randperm(d)`.
 - The value $h(x_i)$ will be different based on the permutation.
- Neat fact:
 - **Probability that $h(x_i) = h(x_j)$ is the Jaccard similarity between x_i and x_j .**
- Proof idea:
 - Probability that you stop with $h(x_i) = h(x_j)$ is given by probability that $x_{ik}=x_{jk}=1$ for a random 'k', divided by probability that at least one of $x_{ik}=1$ or $x_{jk}=1$ is true for a random 'k'.

Low-Memory Randomized Jaccard Approximation

- The “neat fact” lets us approximate Jaccard similarity without storing the shingles.
- First we generate a bunch of random permutations.
 - In practice, use a random hash function to randomly map $1:d$ to $1:d$.
- For each example, go through its shingles to compute $h(x_i)$ for each permutation.
 - No need to store the shingles.
- Approximate $\text{Jaccard}(x_i, x_j)$ as the fraction of permutations where $h(x_i) = h(x_j)$.